

# Multiclass Support Vector Machine exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

In this exercise you will:

- implement a fully-vectorized loss function for the SVM
- implement the fully-vectorized expression for its analytic gradient
- check your implementation using numerical gradients
- use a validation set to tune the learning rate and regularization strength
- optimize the loss function with SGD
- visualize the final learned weights

```
In [5]: # Run some setup code for this notebook.
from __future__ import print_function
import random
import numpy as np
from cs682.data.util import load_CIFAR10
import matplotlib.pyplot as plt

# This is a bit of magic to make matplotlib figures appear inline in the
# notebook rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

## CIFAR-10 Data Loading and Preprocessing

```
In [6]: # Load the raw CIFAR-10 data
cifar10_dir = 'cs682/datasets/cifar-10-batches-py'

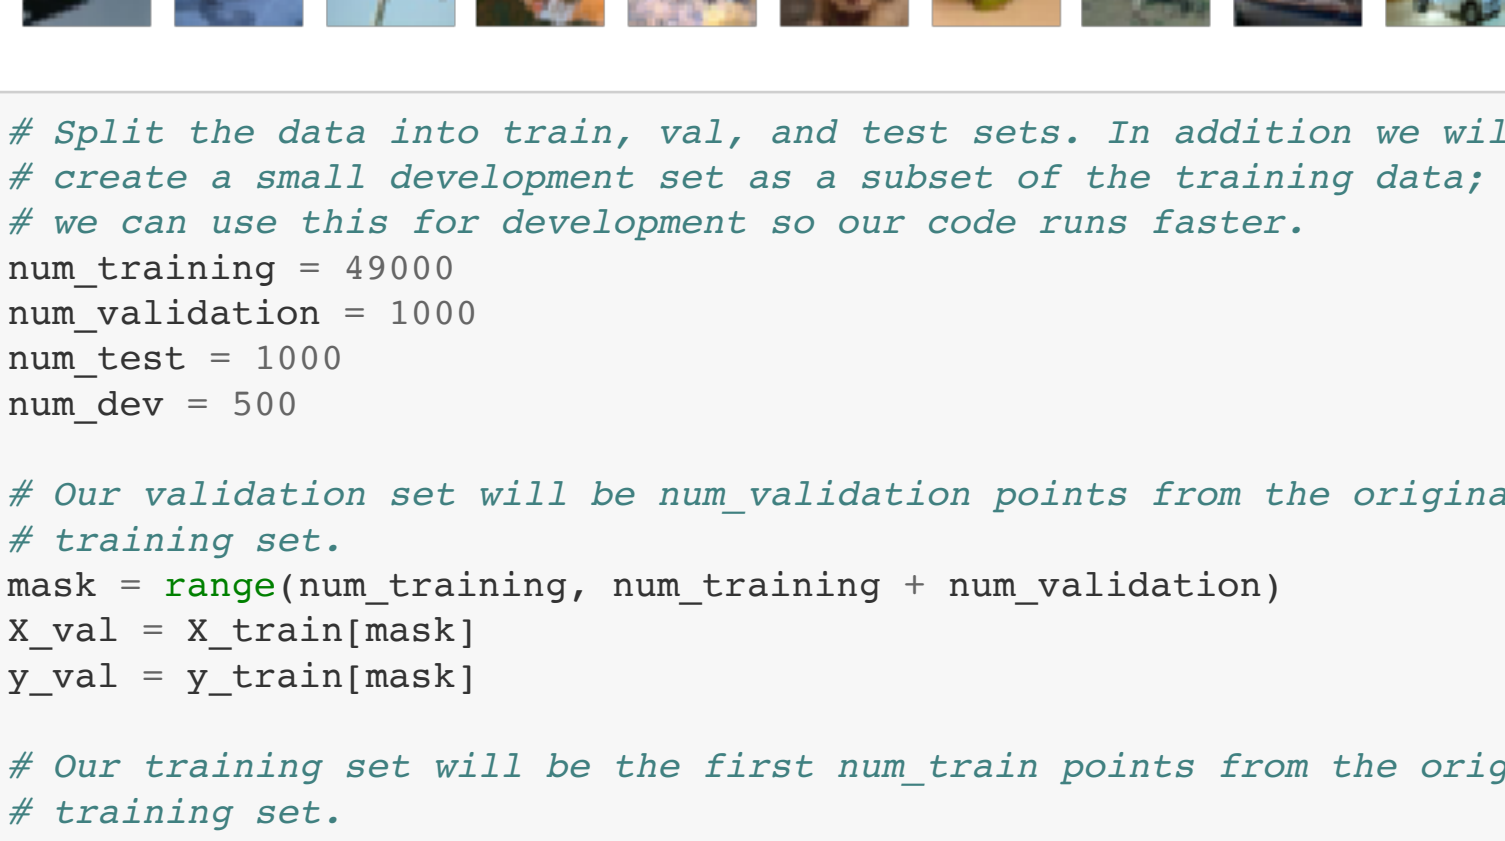
# Cleaning up variables to prevent loading data multiple times (which may cause memory issue)
try:
    del X_train, y_train
    del X_test, y_test
except:
    pass

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

```
In [7]: # Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
num_classes = len(classes)
samples_per_class = 7
for i, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == cls)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```



```
In [8]: # Split the data into train, val, and test sets. In addition we will
# create a small development set as a subset of the training data.
# We can use this set for development so our code runs faster.
num_training = 49000
num_validation = 1000
num_test = 1000
num_dev = 500

# Our validation set will be num_validation points from the original
# training set.
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]

# Our training set will be the first num_train points from the original
# training set.
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

# We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

# We use the first num_test points of the original test set as our
# test set.
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', y_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

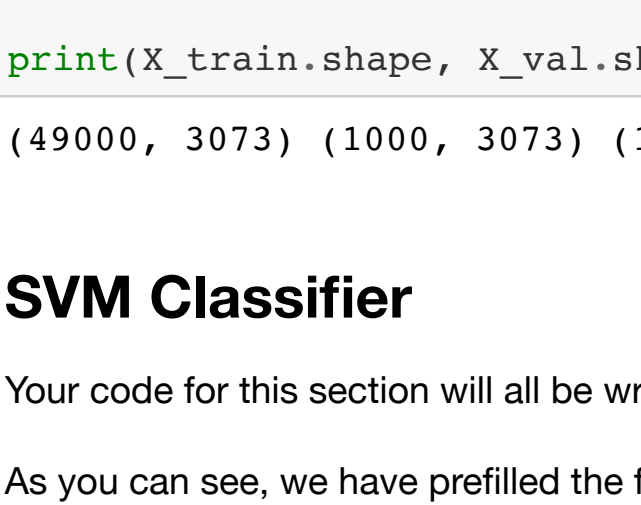
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
```

```
In [9]: # Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('dev data shape: ', X_dev.shape)

Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

```
In [10]: # Preprocessing: subtract the mean image
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)
print(mean_image[0:10]) # print a few of the elements
plt.figure(figsize=(4, 4))
plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean image
plt.show()
```



```
In [11]: # Second: subtract the mean image from train and test data
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image
```

```
In [12]: # third: append the bias dimension of ones (i.e. bias trick) so that our SVM
# only has to worry about optimizing a single weight matrix W.
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)
(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)
```

## SVM Classifier

Your code for this section will all be written inside `cs682/classifiers/linear_svm.py`.

As you can see, we have prefixed the function `svm_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

```
In [13]: # Evaluate the naive implementation of the loss we provided for you:
from cs682.classifiers.linear_svm import svm_loss_naive
import time

# generate a random SVM weight matrix of small numbers
W = np.random.randn(3073, 10) * 0.0001

loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000005)
print('loss: %f' % (loss, ))

loss: 8.606381
```

The grad returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you:

```
In [14]: # Once you've implemented the gradient, recompute it with the code below
# and gradient check it with the function we provided for you

# Compute the loss and its gradient at W.
loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)

# Numerically compute the gradient along several randomly chosen dimensions, and
# compare them with the analytically computed gradient. The numbers should match
# almost exactly along all dimensions.
from cs682.gradient_check import grad_check_sparse
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad)

# do the gradient check once again with regularization turned on
# you didn't forget the regularization gradient did you?
loss, grad = svm_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad)

numerical: 32.662110 analytic: 32.662110, relative error: 1.090752e-11
numerical: 1.663175 analytic: 1.663175, relative error: 3.914669e-11
numerical: -39.167971 analytic: -39.167971, relative error: 9.851151e-12
numerical: 26.852921 analytic: 26.852921, relative error: 3.227127e-12
numerical: -14.226808 analytic: -14.226808, relative error: 1.909513e-11
numerical: 20.146966 analytic: 20.146966, relative error: 3.368709e-12
numerical: -7.677977 analytic: -7.677977, relative error: 5.282179e-11
numerical: -7.721987 analytic: -7.721987, relative error: 3.40845e-11
numerical: 13.315709 analytic: 13.315709, relative error: 3.016378e-12
numerical: -7.827410 analytic: -7.827410, relative error: 4.291792e-11
numerical: 7.752431 analytic: 7.752431, relative error: 1.451516e-11
numerical: 0.760045 analytic: 0.760045, relative error: 4.056273e-10
numerical: -5.437939 analytic: -5.437939, relative error: 5.611144e-11
numerical: 30.272662 analytic: 30.272662, relative error: 8.995795e-12
numerical: 12.690939 analytic: 12.690939, relative error: 2.237610e-11
numerical: -8.214033 analytic: -8.214033, relative error: 5.088870e-11
numerical: -53.618153 analytic: -53.618153, relative error: 6.217904e-12
numerical: -46.852087 analytic: -46.852087, relative error: 4.444220e-12
numerical: -24.371575 analytic: -24.371575, relative error: 1.975149e-12
numerical: 13.691466 analytic: 13.691466, relative error: 1.827882e-11
```

## Inline Question 1:

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? Is it a reason for concern? What is a simple example in one dimension where a gradient check could fail? How would change the margin affect of the frequency of this happening? *Hint: The SVM loss function is not strictly speaking differentiable*

Your Answer: Yes, it is possible that a dimension of grad\_check may be different. Since grad\_check computes the gradient by simply adding a small value to X and checking the function values instead of taking the actual derivative it could be a different value. Also if there is a point where the loss isn't differentiable (like where the hinge loss goes from 0 to above 0) the numerical derivative will return a correct value but the analytical derivative may not.

```
In [15]: # Next implement the function svm_loss_vectorized; for now only compute the loss;
# we will implement the gradient in a moment.
tic = time.time()
loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Naive loss: %e computed in %fs' % (loss_naive, toc - tic))

from cs682.classifiers.linear_svm import svm_loss_vectorized
tic = time.time()
loss_vectorized, _ = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

# The losses should match but your vectorized implementation should be much faster.
print('difference: %f' % (loss_naive - loss_vectorized))

Naive loss: 8.606381e+00 computed in 0.140188s
Vectorized loss: 8.606381e+00 computed in 0.005475s
difference: -0.000000
```

```
In [16]: # Complete the implementation of svm_loss_vectorized, and compute the gradient
# of the loss function in a vectorized way.

# The naive implementation and the vectorized implementation should match, but
# the vectorized version should still be much faster.
tic = time.time()
_, grad_vectorized = svm_loss_naive(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Naive loss and gradient: computed in %fs' % (toc - tic))

tic = time.time()
_, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000005)
toc = time.time()
print('Vectorized loss and gradient: computed in %fs' % (toc - tic))

# The loss is a single number, so it is easy to compare the values computed
# by the two implementations. The gradient on the other hand is a matrix, so
# we use the Frobenius norm to compare them.
difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('difference: %f' % difference)

Naive loss and gradient: computed in 0.104091s
Vectorized loss and gradient: computed in 0.004046s
difference: 0.000000
```

## Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss.

```
In [17]: # In the file linear_classifier.py, implement SGD in the function
# LinearClassifier.train() and then run it with the code below.
from cs682.classifiers import LinearSVM
svm = LinearSVM()
tic = time.time()
loss_hist = svm.train(X_train, y_train, learning_rate=1e-7, reg=2.5e4,
num_iters=1500, verbose=True)
toc = time.time()
print('That took %fs' % (toc - tic))

iteration 0 / 1500: loss 796.140715
iteration 100 / 1500: loss 290.013147
iteration 200 / 1500: loss 109.459323
iteration 300 / 1500: loss 42.935666
iteration 400 / 1500: loss 19.140520
iteration 500 / 1500: loss 10.610751
iteration 600 / 1500: loss 7.287787
iteration 700 / 1500: loss 5.990759
iteration 800 / 1500: loss 5.267446
iteration 900 / 1500: loss 5.252881
iteration 1000 / 1500: loss 4.651852
iteration 1100 / 1500: loss 5.187224
iteration 1200 / 1500: loss 4.735986
iteration 1300 / 1500: loss 6.000257
iteration 1400 / 1500: loss 5.278133
That took 4.907961s
```

```
In [18]: # A useful debugging strategy is to plot the loss as a function of
# training number.
plt.plot(loss_hist)
plt.xlabel('Iteration number')
plt.ylabel('Loss value')
plt.show()
```



```
In [19]: # Write the LinearSVM.predict function and evaluate the performance on both the
# training and validation set.
y_train_pred = svm.predict(X_train)
print('training accuracy: %f' % (np.mean(y_train == y_train_pred), ))
y_val_pred = svm.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred), ))

training accuracy: 0.367776
validation accuracy: 0.377000
```

```
In [21]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rate and regularization strength; if you are careful you should be able to
# get a classification accuracy of about 0.4 on the validation set.
learning_rates = [1e-6, 1e-7]
regularization_strengths = [2.5e4, 1e4]

# results is dictionary mapping tuples of the form
# (learning_rate, regularization_strength) to tuples of the form
# (training_accuracy, validation_accuracy). The accuracy is simply the fraction
# of data points that are correctly classified.
results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation rate.

#####
# TODO:
# Write code that chooses the best hyperparameters by tuning on the validation
# set. For each combination of hyperparameters, train a Linear SVM on the
# training set, compute its accuracy on the training and validation sets, and
# store these numbers in the results dictionary. In addition, store the best
# validation accuracy, best_val and the LinearSVM object that achieves this
# accuracy in best_svm.
#
# Hint: You should use a small value for num_iters as you develop your
# validation code so that the SVMs don't take much time to train; once you are
# confident that your validation code works, you should run the validation
# code with a larger value for num_iters.
#####

for lr in learning_rates:
    for reg in regularization_strengths:
        print("Testing LR=%f, reg=%f" % (lr, reg))
        svm = LinearSVM()
        svm.train(X_train, y_train, learning_rate=lr, reg=reg,
num_iters=10000, verbose=False)

        y_train_pred = svm.predict(X_train)
        y_val_pred = svm.predict(X_val)
        val_acc = np.mean(y_val == y_val_pred)
        print('validation accuracy: %f' % val_acc)

        results[(lr, reg)] = (train_acc, val_acc, svm)

#####
# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy, svm = results[(lr, reg)]
    if val_accuracy > best_val:
        best_val = val_accuracy
        best_svm = svm
        print('lr %f reg %f train accuracy: %f val accuracy: %f' % (
            lr, reg, train_accuracy, val_accuracy))


print('best validation accuracy achieved during cross-validation: %f' % best_val)
```

```
Testing LR= 2e-08 reg= 25000.0
training accuracy: 0.374898
validation accuracy: 0.394000
Testing LR= 2e-08 reg= 10000.0
training accuracy: 0.385992
validation accuracy: 0.403000
Testing LR= 1e-07 reg= 25000.0
training accuracy: 0.387000
validation accuracy: 0.373000
Testing LR= 1e-07 reg= 10000.0
training accuracy: 0.384004
validation accuracy: 0.382000
lr 2.000000e-08 reg 1.000000e+04 train accuracy: 0.389592 val accuracy: 0.403000
lr 2.000000e-08 reg 1.000000e+04 train accuracy: 0.374898 val accuracy: 0.394000
lr 1.000000e-07 reg 1.000000e+04 train accuracy: 0.384204 val accuracy: 0.382000
lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.367000 val accuracy: 0.373000
validation accuracy achieved during cross-validation: 0.403000
```

```
In [22]: # Visualize the cross-validation results
import math
x_scatter = [math.log10(x[i]) for x in results]
y_scatter = [math.log10(y[i]) for x in results]

# plot training accuracy
marker_size = 100
colors = [results[x][0] for x in results]
plt.subplot(2, 1, 1)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 training accuracy')

# plot validation accuracy
colors = [results[x][1] for x in results] # default size of markers is 20
plt.subplot(2, 1, 2)
plt.scatter(x_scatter, y_scatter, marker_size, c=colors)
plt.colorbar()
plt.xlabel('log learning rate')
plt.ylabel('log regularization strength')
plt.title('CIFAR-10 validation accuracy')
plt.show()
```

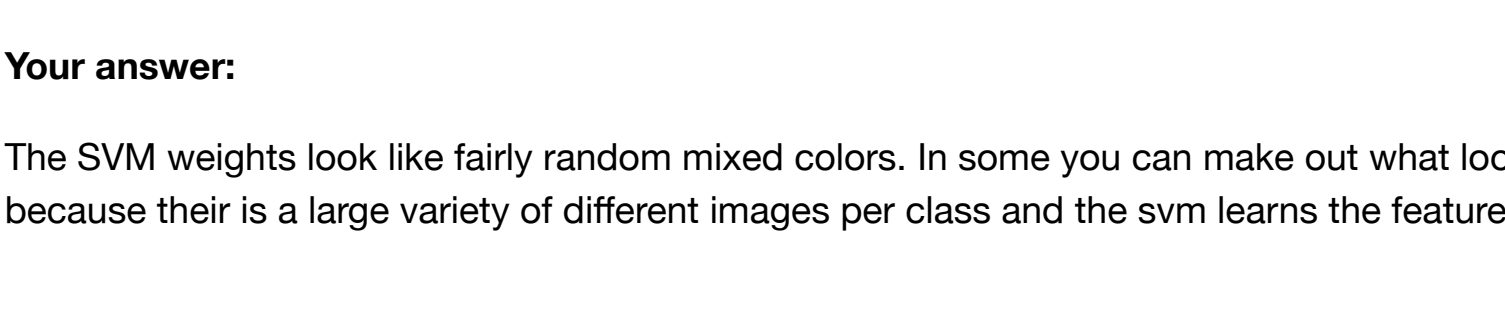


```
In [24]: # Evaluate the best svm on test set
y_test_pred = best_svm.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)

linear SVM on raw pixels final test set accuracy: 0.382000
```

```
In [25]: # Visualize the learned weights for each class.
# Depending on your choice of learning rate and regularization strength, these may
# or may not be nice to look at.
w = best_svm.W[1:-1, 1:] # strip out the bias
w = w.reshape(32, 32, 3, 10)
w_min, w_max = np.min(w, axis=(0, 1)), np.max(w, axis=(0, 1))
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[i, :, :, :].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```



## Inline question 2:

Describe what your visualized SVM weights look like, and offer a brief explanation for why they look they way that they do.

Your answer: The SVM weights look like fairly random mixed colors. In some you can make out what looks like images like the truck, car, horse and frog. They look so mixed because there is a large variety of different images per class and the svm learns the features most common to each class.